

# Transactions and Concurrency

---

EGCI 321: LECTURE 14 (WEEK 9)

# Outline

---

## 1. Why We Need Transactions

- Failures
- Concurrency

## 2. Serializability

- Serialization Schedules
- Serialization Graphs

## 3. Transactions in SQL

- Abort and Commit
- Isolation Levels

## 4. Implementing Transactions

- Concurrency Control
- Recovery Management

# Problems Caused by Failures

---

- Update all account balance at a bank branch  
`Accounts(Anum, Cid, BranchId, Balance)`

**update**      `Accounts`  
**set**          `Balance = Balance * 1.05`  
**where**        `BranchId = 12345`

## Problem

If the system crashes while processing this update, some, but, not all, tuples with **BranchID = 12345** may have been updated

# Another Failure-Related Problem

---

- Transfer money between accounts:

**update** Accounts

**set**                     $\text{Balance} = \text{Balance} - 100$

**where**  $\text{Anum} = 8888$

**update** Accounts

**set**                     $\text{Balance} = \text{Balance} + 100$

**where**  $\text{Anum} = 9999$

## **Problem**

If the system fails between these updates, money may be withdrawn but not redeposited.

# Problems Caused by Concurrency

---

## Applications 1:

**update** Accounts

**set** Balance = Balance – 100

**where** Anum = 8888

**update** Accounts

**set** Balance = Balance + 100

**where** Anum = 9999

## Application 2:

**select** Sum(Balance)

**from** Accounts

## Problem

If the applications run concurrently, the total balance returned to application 2 may be inaccurate.

# Another Concurrency Problem

---

## Application 1:

```
select          balance into : balance
from            Accounts
where   Anum = 8888

compute        :newbalance using :balance

update Accounts
set            Balance = :newbalance
where   Anum = 8888
```

## Application 2: same as Application 1

## Problem

If the applications run concurrently, one of the updates may be “lost”

# Transactions

---

## Definition (Transaction)

An application-specified *atomic* and *durable* unit of work.

## Properties of transactions ensured by the DBMS:

**Atomic:** a transaction occurs entirely, or not at all

**Consistency:** each transaction preserves the consistency of the database

**Isolated:** concurrent transactions do not interfere with each other

**Durables:** once completed, a transaction's changes are permanent

# Serializability (informal)

---

Concurrent transactions must appear to have been executed sequentially, i.e., one at a time, in some order. If  $T_i$  and  $T_j$  are concurrent transactions, then either:

1.  $T_i$  will appear to precede  $T_j$ , meaning that  $T_j$ , will “see” any updates made by  $T_i$ , and  $T_i$  will not see any updates made by  $T_j$ , or
2.  $T_i$  will appear to follow  $T_j$ , meaning that  $T_i$  will see  $T_j$ ’s updates and  $T_j$  will not see  $T_i$ ’s



# Serializability: An Example

---

An interleaved execution of two transactions,  $T_1$  and  $T_2$ :

$$H_a = w_1[x] r_2[x] w_1[y] r_2[y]$$

An equivalent serial execution of  $T_1$  and  $T_2$

$$H_b = w_1[x] w_1[y] r_2[x] r_2[y]$$

An interleaved execution of  $T_1$  and  $T_2$  with no equivalent serial execution:

$$H_c = w_1[x] r_2[x] r_2[y] w_1[y]$$

$H_a$  is serializable because it is equivalent to  $H_b$ , a serial schedule  $H_c$  is not serializable

# Transactions and Histories

---

- Two operations conflict if:
  1. They belong to different transactions,
  2. They operate on the same object, and
  3. At least one of the operations is a write
- A transaction is a sequence of read and write operations
- A *execution history* over a set of transactions  $T_1 \dots T_n$  is an interleaving of the operations of  $T_1 \dots T_n$  in which the operation ordering imposed by each transaction is preserved
- Two important assumptions:
  1. Transactions interact with each other only via reads and writes of objects
  2. A database is a fixed set of independent objects

# Serializability

---

## Definition ((Conflict) Equivalence)

Two histories are *(conflict) equivalent* if

- They are over the same set of transactions, and
- The ordering of each pair of conflicting operations is the same in each history

## Definition ((Conflict) Serializability)

A history  $H$  is said to be *(conflict) serializable* if there exists some serial history  $H'$  that is (conflict) equivalent to  $H$

# Testing for Serializability

---

$r_1[x] \ r_3[x] \ w_4[y] \ r_2[u] \ w_4[z] \ r_1[y] \ r_3[u] \ r_2[z] \ w_2[z] \ r_3[z] \ r_1[z] \ w_3[y]$


Is this history serializable?

## Theorem

A history is serializable iff its serialization graph is acyclic

# Serialization Graphs

$r_1[x]$   $r_3[x]$   $w_4[y]$   $r_2[u]$   $w_4[z]$   $r_1[y]$   $r_3[u]$   $r_2[z]$   $w_2[z]$   $r_3[z]$   $r_1[z]$   $w_3[y]$



Consider read-write, write-read, and write-write of same object but different transactions  
Ex.  $w_4, r_1, w_3, \dots$

*Conflict pairs*  $y: (w_4[y] r_1[y]), (r_1[y] w_3[y]), (w_4[y] w_3[y])$


# Serialization Graph (cont.)

$r_1[x]$   $r_3[x]$   $w_4[y]$   $r_2[u]$   $w_4[z]$   $r_1[y]$   $r_3[u]$   $r_2[z]$   $w_2[z]$   $r_3[z]$   $r_1[z]$   $w_3[y]$

Conflict pairs

$y: (w_4[y] r_1[y]), (r_1[y] w_3[y]), (w_4[y] w_3[y])$

$z: (w_4[z] r_2[z]), (w_2[z] r_3[z]), (w_2[z] r_1[z])$



The history above is equivalent to

$w_4[y]$   $w_4[z]$   $r_2[u]$   $r_2[z]$   $w_2[z]$   $r_1[x]$   $r_1[y]$   $r_1[z]$   $r_3[x]$   $r_3[u]$   $r_3[z]$   $w_3[y]$

That is, it is equivalent to executing  $T_4$  followed by  $T_2$  followed by  $T_1$  followed by  $T_3$

# Abort and Commit

---

- A transaction may terminate in one of two ways:
  - When a transaction **commits**, any updates it made become durable, and they become visible to other transactions. A commit is the “all” in “all-or-nothing” execution
  - When a transaction **aborts**, any updates it may have made are undone (erased), as if the transaction never ran at all. An abort is the “nothing” in “all-or-nothing” execution
- A transaction that has started but has not yet aborted or committed is said to be **active**

# Transactions in SQL

---

- A new transaction is begun when an application first executes an SQL command
- Two SQL commands are available to terminate a transaction:
  - **commit work:** commits the transaction
  - **rollback work:** abort the transaction
- A new transaction begins with the application's next SQL command after **commit work** or **rollback work**



# SQL Isolation Levels

---

- SQL allows the serializability guarantee to be relaxed, if necessary
- For each transaction, it is possible to specify an *isolation level*
- For isolation levels are supported, with the highest being serializability:

**Level 0 (Read Uncommitted):** transaction may see uncommitted updates

**Level 1 (Read Committed):** transaction sees only committed changes, but non-repeatable reads are possible

**Level 2 (Repeatable Read):** reads are repeatable, but “phantoms” are possible

**Level 3 (Serializability)**

# Non-Repeatable Reads

---

## Application 1:

```
update Employee  
set           Salary = Salary +1000  
where WorkDept = 'D11'
```

## Application 2:

```
select        * from Employee  
where WorkDept = 'D11'  
select        * from Employee  
where Lastname Like 'A%'
```

## Problem

If there are employees in D11 with surnames that begin with “A”, Application 2’s queries may see them with different salaries

# Phantoms

---

## Application 1:

```
insert into      Employee
values ( '000123', 'Sheldon', 'Q', 'Jetstream', 'D11', "05/01/00, 520000.00)
```

## Application 2

```
select  *
from    Employee
where   WorkDept = 'D11'
select  *
from    Employee
where   Salary > 5000
```

## Problem

Application 2's second query may see Sheldon Jetstream, even though its first query does not

# Implementing Transactions

---

The implementation of transactions in a DBMS has two parts:

**Concurrency Control:** guarantees that the execution history has the desired properties ( such as serializability )

**Recovery Management:** guarantee that committed transactions are durable (despite failures), and that aborted transactions have no effect on the database

# Concurrency Control

---

- Serializability can be guaranteed by executing transactions serially, but, in many environments, this leads to poor performance
- Typically, many transactions are in progress concurrently, and a concurrency control protocol is used to ensure that the resulting history is serializable
- Many concurrency control protocols have been proposed, based on:
  - Locking, or
  - Timestamps, or
  - Serialization graph analysis
- By far the most commonly implemented protocol is *strict two-phase locking*
- The strict two-phase locking protocol can be relaxed, as necessary, to accommodate isolation levels below serializability

# Strict Two-Phase Locking

---

## The rules:

1. Before a transaction may read or write an object, it must have a lock on the object
  - A shared lock is required to read an object
  - An exclusive lock is required to write an object
2. Two or more transactions may not hold locks on the same object unless all hold shared locks
3. A transaction may not release any locks until it commits (or aborts)

If all transactions use strict two-phase locking, the execution history is guaranteed to be serializable

# Transaction Blocking

---

- Consider the following sequence of events:
  - $T_1$  acquires a shared lock on  $x$  and reads  $x$
  - $T_2$  attempts to acquire an exclusive lock on  $x$  (so that it can write  $x$ )
- The two-phase locking rules prevent  $T_2$  from acquiring its exclusive lock—this is called a *lock conflict*
- Lock conflicts can be resolved in one of two ways:
  1.  $T_2$  can be *blocked* – forced to wait until  $T_1$  release its lock
  2.  $T_1$  can be *pre-empted* – forced to abort and give up its locks

# Deadlocks

---

- Transaction blocking can result in deadlocks
- For example:
  - $T_1$  reads object  $x$
  - $T_2$  reads object  $y$
  - $T_2$  attempts to write object  $x$  (it is blocked)
  - $T_1$  attempts to write object  $y$  (it is blocked)

A deadlock can be resolved only by forcing one of the transactions involved in the deadlock to abort



# Recovery Management

---

*Recovery management* means:

1. Implementing voluntary or involuntary rollback of individual transactions
2. Implementing recovery from *system failures*

*System failure* mean:

1. The database server is halted abruptly
2. Processing of in-progress SQL command(s) is halted abruptly
3. Connections to Applications programs (clients) are broken.
4. Contents of memory buffers are lost
5. Database files are damaged

# Failures and Transactions

---

- To ensure that transactions are atomic, every transaction that is active when a system failure occurs must either be
  - Restarted after the failure from the point it which it left off, or
  - Rolled back after the failure
- It is difficult to restart applications after a system failure, so the recovery manager does the following:
  - Abort transactions that were active at the time of the failure
  - Ensure that changes made by transactions that committed before the failure are not lost

# Logging

---

- Recovery management is usually accomplished using a *log*
- A log is a read/append data structure located in persistent storage (it must survive the failure)
- When transactions are running, *log records* are appended to the log. Log records contain:
  - UNDO information:** old version of objects that have been modified by a transaction. Used to undo database changes made by a transaction that aborts.
  - REDO information:** new versions of objects that have been modified by a transaction. Used to redo the work done by a transaction that commits.
  - BEGIN/COMMIT/ABORT:** records are recorded whenever a transaction begins, commits, or aborts.

## Requires Write-Ahead-Logging

Log records must be written *before* updating the database

# Recovery

---

## Recovering from a system failure:

### 1. Scan the log from tail (newest) to head (oldest):

- ▶ Create a list of committed transactions
- ▶ Undo updates of active and aborted transactions

### 2. Scan the log from head (oldest) to tail (newest):

- ▶ Redo updates of committed transactions

## Rolling back a single transaction:

### 1. Scan the log from the tail to the transaction's BEGIN record.

- ▶ Undo the transaction's updates

# Reference

---

1. Ramakrishnan R, Gehrke J., Database management systems, 3<sup>rd</sup> ed., New York (NY): McGraw-Hill, 2003.