



# System Programming

## Module 8



# EGCI 252

# System Programming

Computer Engineering Department  
Faculty of Engineering



# Shared Memory

- Normally, the Unix kernel prohibits one process from accessing (reading, writing) memory belonging to another process
- Sometimes, however, this restriction is inconvenient
- At such times, System V IPC Shared Memory can be created to specifically allow one process to read and/or write to memory created by another process



# Advantages of Shared Memory

- Random Access
  - you can update a small piece in the middle of a data structure, rather than the entire structure
- Efficiency
  - unlike message queues and pipes, which copy data from the process *into* memory within the kernel, shared memory is directly accessed
  - Shared memory resides in the user process memory, and is then shared among other processes



# Disadvantages of Shared Memory

- No automatic synchronization as in pipes or message queues (you have to provide any synchronization). Synchronize with *semaphores* or signals.
- You must remember that pointers are only valid within a given process. Thus, pointer offsets cannot be assumed to be valid across inter-process boundaries. This complicates the sharing of linked lists or binary trees.



# Creating Shared Memory

```
int shmget(key_t key, size_t size, int shmflg);
```

- key is either a number or the constant IPC\_PRIVATE (man ftok)
- a shmid is returned
- key\_t ftok(const char \* path, int id) will return a key value for IPC usage
- size is the size of the shared memory data
- shmflg is a rights mask (0666) OR'd with one of the following:
  - IPC\_CREAT will create or attach
  - IPC\_EXCL creates new or it will error if it exists



# Attaching Shared Memory

- After obtaining a shmid from shmget(), you need to *attach* or map the shared memory segment to the address space of a process:

```
void * shmat(int shmid, void * shmaddr, int shmflg)
```

- shmid is the id returned from shmget()
- shmaddr is the shared memory segment address. Set this to NULL and let the system handle it.
- shmflg is one of the following (usually 0):
  - SHM\_RDONLY sets the segment readonly
  - SHM\_RND sets page boundary access
  - SHM\_SHARE\_MMU set first available aligned address



# Detaching Shared Memory

- In order to detach the shared memory from the current process, the shmdt function must be issued.

`void * shmdt(void * first_byte_addr)`

- `first_byte_addr` is a pointer to the address returned by `shmat()` function.



# Shared Memory Control

```
struct shmid_ds {  
    int shm_segsz;      /* size of segment in bytes */  
    __time_t shm_atime; /* time of last shmat command */  
    __time_t shm_dtime; /* time of last shmdt command */  
    ...  
    unsigned short int __shm_npages; /* size of segment in pages */  
    msgqnum_t shm_nattach; /* number of current attaches */  
    ... /* pids of creator and last shmop */  
};
```

- `int shmctl(int shmid, int cmd, struct shmid_ds * buf);`
- `cmd` can be one of:
  - `IPC_RMID` destroy the memory specified by `shmid`
  - `IPC_SET` set the `uid`, `gid`, and mode of the shared mem
  - `IPC_STAT` get the current `shmid_ds` struct for the queue



# Shared Memory - Example

```
//Consumer.c
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define MEM_SIZE 4096

struct shm_st
{
    int written;
    char data[BUFSIZ];
}

int main()
{
    int running = 1, shmid;
    void *sh_mem = NULL;
    struct shm_st *sh_area;
```

```
    srand((unsigned int) getpid());
    shmid = shmget((key_t) 1234, MEM_SIZE,
                  0666 | IPC_CREAT);

    if (shmid == -1)
    {
        fprintf(stderr, "shmget failed\n");
        exit(EXIT_FAILURE);
    }

    sh_mem = shmat(shmid, NULL, 0);
    if (sh_mem == (void *) -1)
    {
        fprintf(stderr, "shmat failed\n");
        exit(EXIT_FAILURE);
    }

    printf("Memory attached at %X\n", sh_mem);
    sh_area = (struct shm_st *) sh_mem;
    sh_area->written = 0;
```



# Shared Memory - Example (Cont.)

```
while (running)
{
    if (sh_area->written)
    {
        printf("Data written in shared memory: %s\n", sh_area-> data);
        sh_area->written = 0;
        if (strncmp(sh_area->data, "end", 3) == 0)
            running = 0;
    }
    sleep(rand() % 4);
}

if (shmdt(sh_mem) == -1 || shmctl(shmid, IPC_RMID, 0) == -1)
{
    fprintf(stderr, "shmdt or shmctl failed\n");
    exit(EXIT_FAILURE);
}

exit(EXIT_SUCCESS);
}
```



# Shared Memory – Example (Cont.)

```
//Producer.c
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define MEM_SIZE 4096

struct shm_st
{
    int written;
    char data[BUFSIZ];
}

int main()
{
    int running = 1, shmID;
    void *sh_mem = NULL;
    struct shm_st *sh_area;
```

```
    char buffer[BUFSIZ];

    shmID = shmget((key_t) 1234, MEM_SIZE,
                  0666 | IPC_CREAT);
    if (shmID == -1)
    {
        fprintf(stderr, "shmget failed\n");
        exit(EXIT_FAILURE);
    }

    sh_mem = shmat(shmID, NULL, 0);
    if (sh_mem == (void *) -1)
    {
        fprintf(stderr, "shmat failed\n");
        exit(EXIT_FAILURE);
    }

    printf("Memory attached at %X\n", sh_mem);
    sh_area = (struct shm_st *) sh_mem;
```



# Shared Memory – Example (Cont.)

```
while (running)
{
    while (sh_area->written)
    {
        sleep(1); printf("Waiting...\n");
    }
    printf("Enter data: ");
    fgets(buffer, BUFSIZ, stdin);
    strcpy(sh_area->data, buffer);
    sh_area->written = 1;
    if (strncmp(buffer, "end", 3) == 0)
        running = 0;
}

if (shmdt(sh_mem) == -1)
{
    fprintf(stderr, "shmdt failed\n");
    exit(EXIT_FAILURE);
}
exit(EXIT_SUCCESS);
}
```



# Assignment

- Write a simple chat program using a shared memory
- Requirements:
  - The program's name must be "schat.c"
  - The program takes one command line argument (i.e., 1 or 2 to indicate the type of messages)
  - Create a shared memory with the key value of 21930
  - Must be able to concurrently send and receive any messages between two "schat" processes by using the shared memory.
  - Use the word "end chat" as a command to end the chat program



**End of Module**